

Latent Variable Models


Volodymyr Kuleshov

Cornell Tech


Lecture 5

The Task of Generative Modeling

Suppose we are given a training set of examples, e.g., images of dogs



$$\begin{aligned} x_i &\sim P_{\text{data}} \\ i &= 1, 2, \dots, n \end{aligned}$$



Our Goal: define a probability distribution $p(x)$ over images x such that

- **Generation:** If we sample $x_{\text{new}} \sim p(x)$, x_{new} should look like a dog.
- **Representation Learning:** We should be able to learn what these images have in common, e.g., ears, tail, etc.
- **Density Estimation:** $p(x)$ should be high if x looks like a dog, and low otherwise (*anomaly detection*)

Previously: Generation and density estimation. Today: Representation learning.

Recap of Previous Lectures

1 Autoregressive models:

- Probability distributions factorize into a product of factors:

$$p(\mathbf{x}) = \prod_{i=1}^n p(x_i \mid \mathbf{x}_{<i})$$

- We can efficiently represent p via *conditional independence* or via compact *neural parameterizations* $p_{\text{Neural}}(x_i \mid \mathbf{x}_{<i})$ of each factor.

2 Pros of autoregressive models:

- It is computationally tractable to evaluate likelihoods
- It is tractable to train $p(\mathbf{x})$ via maximum likelihood & gradient descent

3 Cons of autoregressive models:

- They require choosing an ordering over variables
- Generation is sequential (hence usually slow)
- Cannot learn features in an unsupervised way

- ① Latent Variable Models
 - Motivation
 - Definition
- ② Examples of Shallow and Deep Latent Variable Models
 - Gaussian Mixture Models
 - Deep Latent Gaussian Models
- ③ Approximating Marginal Likelihood Using Variational Inference
 - Challenges of Maximum Marginal Likelihood Learning
 - Evidence Lower Bound on the Marginal Likelihood
 - Variational Inference

Latent Variable Models: Motivation

Human faces \mathbf{x} can feature a lot of interesting characteristics: gender, age, hair color, eye color, pose, etc.



Challenge: How to automatically learn these characteristics from data?

- Unless the images are annotated, these factors of variation are not explicitly available (latent).
- **Idea:** Explicitly model these factors using latent variables \mathbf{z} and learn them using unsupervised learning.

Latent Variable Models: Definition



A latent variable model defines a probability distribution


$$p(x, z) = p(x|z)p(z)$$

containing two sets of variables:

- 1 Observed variables \mathbf{x} that represent the high-dimensional objects we are trying to model and that are in our training set.
- 2 Latent variables \mathbf{z} that are not in the dataset, but that are associated with \mathbf{x} as specified by $p(\mathbf{x}, \mathbf{z})$. We will learn \mathbf{z} and $p(\mathbf{x}, \mathbf{z})$ jointly.

Latent Variable Models: Example


Consider the following distribution $p(\mathbf{x}|\mathbf{z})p(\mathbf{z})$ over faces:



- The observed variables \mathbf{x} are the images (their pixel values)
- The latent variables \mathbf{z} are discrete features: gender, pose, etc.
 - Note that we can extract features via $p(\mathbf{z} | \mathbf{x})$, e.g., $p(\text{Gender} = F | \mathbf{x})$
- **Challenges:**
 - 1 Very difficult to specify conditional $p(\mathbf{x}|\mathbf{z})$ by hand
 - 2 Unsupervised learning in this model can be intractable

Continuous Latent Variable Representations

Our solution to Challenge #1 involves choosing a continuous (“distributed”) representation for \mathbf{z} :




- 1 The variable $\mathbf{z} \in \mathbb{R}^p$ is a continuous real-valued vector
- 2 The $\mathbf{z} \rightarrow \mathbf{x}$ mapping is specified by a neural net and learned from data

The \mathbf{z} form a smooth parameterization of faces \mathbf{x} . Similar faces (same age, gender) have similar \mathbf{z} . We can map \mathbf{x} to \mathbf{z} , interpolate between \mathbf{z} , generate faces from new \mathbf{z} , etc.

Deep Latent Variable Models: Example


Here is an example of how a deep latent variable model with a continuous latent representation \mathbf{z} might look like:



- 1 The prior $p(\mathbf{z})$ is a p -dimensional Gaussian $\mathcal{N}(0, I_p)$
- 2 $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mu_{\theta}(\mathbf{z}), \Sigma_{\theta}(\mathbf{z}))$ where $\mu_{\theta}, \Sigma_{\theta}$ are neural networks


Learning such models is still challenging.

Example: Unsupervised Learning Over Handwritten Digits



Unsupervised clustering of MNIST digits using deep latent variable models.

Example: Unsupervised Learning over Face Images




Benefits of the Latent Variable Approach

- Allows us to define complex models $p(\mathbf{x})$ in terms of simple building blocks $p(\mathbf{x} | \mathbf{z})$
- Natural for unsupervised learning tasks (clustering, unsupervised representation learning, etc.)
- No free lunch: much more difficult to learn compared to fully observed, autoregressive models

- ① Latent Variable Models
 - Motivation
 - Definition
- ② **Examples of Shallow and Deep Latent Variable Models**
 - Gaussian Mixture Models
 - Deep Latent Gaussian Models
- ③ Approximating Marginal Likelihood Using Variational Inference
 - Challenges of Maximum Marginal Likelihood Learning
 - Evidence Lower Bound on the Marginal Likelihood
 - Variational Inference

Shallow Latent-Variable Models: Gaussian Mixture Models

The classical latent variable model $p(\mathbf{x}, z)$ is the mixture of Gaussians:




- 1 The prior is $p(z) = \text{Categorical}(1, \dots, K)$
- 2 The conditional is $p(\mathbf{x} \mid z = k) = \mathcal{N}(\mu_k, \Sigma_k)$

This assumes the data is a mixture of K unknown Gaussian clusters.

Representational Power of Mixture Models

Mixtures of Gaussians can represent distributions that single Gaussians cannot, e.g.: the two-component structure of this data:




Mixtures of Gaussians can be used for generative tasks:

- **Generation:** First sample z , then sample z -th Gaussian to get \mathbf{x}
- **Representation Learning:** Learn cluster components from unlabeled data. The posterior $p(z | \mathbf{x})$ identifies the latent cluster.

Deep Latent Gaussian Models

We can extend GMMs to mixtures of an infinite # of Gaussians:




- ① The prior is $p(\mathbf{z}) = \mathcal{N}(\mathbf{0}, I)$
- ② $p(\mathbf{x} | \mathbf{z}) = \mathcal{N}(\mu_{\theta}(\mathbf{z}), \Sigma_{\theta}(\mathbf{z}))$ where $\mu_{\theta}, \Sigma_{\theta}$ are neural networks; e.g.:
 - $\mu_{\theta}(\mathbf{z}) = \sigma(A\mathbf{z} + c) = (\sigma(a_1\mathbf{z} + c_1), \sigma(a_2\mathbf{z} + c_2)) = (\mu_1(\mathbf{z}), \mu_2(\mathbf{z}))$
 - $\Sigma_{\theta}(\mathbf{z}) = \text{diag}(\exp(\sigma(B\mathbf{z} + d))) = \begin{pmatrix} \exp(\sigma(b_1\mathbf{z} + d_1)) & 0 \\ 0 & \exp(\sigma(b_2\mathbf{z} + d_2)) \end{pmatrix}$
 - $\theta = (A, B, c, d)$

Even though $p(\mathbf{x} | \mathbf{z})$ is simple, the marginal $p(\mathbf{x})$ is very complex/flexible

Maximum Marginal Likelihood for Latent Variable Models

We may learn latent variable models by choosing p that maximizes the marginal likelihood $p(\mathbf{x})$ averaged over \mathbf{x} . For Gaussian mixtures, we have:

$$p(\mathbf{x}) = \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z}) = \sum_{\mathbf{z}} p(\mathbf{z}) p(\mathbf{x} | \mathbf{z}) = \sum_{k=1}^K p(\mathbf{z} = k) \underbrace{\mathcal{N}(\mathbf{x}; \mu_k, \Sigma_k)}_{k\text{-th component}}$$



These objectives are non-convex, and in the case of VAEs they're intractable to compute, which motivates approximations we will see next.

- ① Latent Variable Models
 - Motivation
 - Definition
- ② Examples of Shallow and Deep Latent Variable Models
 - Gaussian Mixture Models
 - Deep Latent Gaussian Models
- ③ **Approximating Marginal Likelihood Using Variational Inference**
 - Challenges of Maximum Marginal Likelihood Learning
 - Evidence Lower Bound on the Marginal Likelihood
 - Variational Inference

Running Example: A Generative Model for MNIST

Suppose we are given a dataset \mathcal{D} of handwritten digits (binarized MNIST)



- Each image has $n = 28 \times 28 = 784$ pixels. Each pixel can either be black (0) or white (1).
- **Our Goal:** Learn a probability distribution $p(x) = p(x_1, \dots, x_{784})$ over $x \in \{0, 1\}^{784}$ such that when $x \sim p(x)$, x looks like a digit.

Running Example: MNIST with Missing Values

One way to introduce latent variables into our running example is to assume that some pixels are missing at train time (e.g., top half):



- Let \mathbf{x} denote the observed random variables, and \mathbf{z} the unobserved ones (also called hidden or latent)
- Suppose we have a model (e.g., an autoregressive model) for the joint

$$p(\mathbf{x}, \mathbf{z}; \theta).$$

What is the marginal likelihood $p(\mathbf{x} = \bar{\mathbf{x}}; \theta)$ of a training data point $\bar{\mathbf{x}}$?

$$\sum_{\bar{\mathbf{z}}} p(\mathbf{x} = \bar{\mathbf{x}}, \mathbf{z} = \bar{\mathbf{z}}; \theta) = \sum_{\mathbf{z}} p(\bar{\mathbf{x}}, \bar{\mathbf{z}}; \theta)$$

This sums the probabilities of all possible completions of the image (green)

Challenges of Maximum Marginal Likelihood Learning

- Suppose we have a dataset \mathcal{D} , where for each datapoint the \mathbf{x} variables are observed (e.g., pixel values) and the variables \mathbf{z} are never observed (e.g., cluster or class id.). $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M)}\}$.
- Recall that maximum likelihood learning involves maximizing:

$$\log \prod_{\mathbf{x} \in \mathcal{D}} p(\mathbf{x}; \theta) = \sum_{\mathbf{x} \in \mathcal{D}} \log p(\mathbf{x}; \theta) = \sum_{\mathbf{x} \in \mathcal{D}} \log \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z}; \theta)$$

- Evaluating $\log \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z}; \theta)$ can be intractable. Suppose we have 30 binary latent features, $\mathbf{z} \in \{0, 1\}^{30}$. Evaluating $\sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z}; \theta)$ involves a sum with 2^{30} terms. For continuous variables, $\log \int_{\mathbf{z}} p(\mathbf{x}, \mathbf{z}; \theta) d\mathbf{z}$ is often intractable. Gradients ∇_{θ} also hard to compute.
- Need **approximations**. One gradient evaluation per training data point $\mathbf{x} \in \mathcal{D}$, so approximation needs to be cheap.

Naive Approach: Naive Monte Carlo

The marginal likelihood $p_{\theta}(\mathbf{x})$ for partially observed data is:

$$p_{\theta}(\mathbf{x}) = \sum_{\text{All values of } \mathbf{z}} p_{\theta}(\mathbf{x}, \mathbf{z}) = |\mathcal{Z}| \sum_{\mathbf{z} \in \mathcal{Z}} \frac{1}{|\mathcal{Z}|} p_{\theta}(\mathbf{x}, \mathbf{z}) = |\mathcal{Z}| \mathbb{E}_{\mathbf{z} \sim \text{Uniform}(\mathcal{Z})} [p_{\theta}(\mathbf{x}, \mathbf{z})]$$

We can think of it as an (intractable) expectation. Monte Carlo to the rescue:

- 1 Sample $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(k)}$ uniformly at random
- 2 Approximate expectation with sample average

$$\sum_{\mathbf{z}} p_{\theta}(\mathbf{x}, \mathbf{z}) \approx |\mathcal{Z}| \frac{1}{k} \sum_{j=1}^k p_{\theta}(\mathbf{x}, \mathbf{z}^{(j)})$$

Works in theory but not in practice. For most \mathbf{z} , $p_{\theta}(\mathbf{x}, \mathbf{z})$ is very low (most completions don't make sense). Some are very large but will never "hit" likely completions by uniform random sampling, hence we will mis-estimate our objective. We need a clever way to select $\mathbf{z}^{(j)}$.

Second Approach: Importance Sampling

Importance sampling is a clever way of applying Monte Carlo. First, we express $p_\theta(\mathbf{x})$ as a different kind of expectation:

$$p_\theta(\mathbf{x}) = \sum_{\mathbf{z} \in \mathcal{Z}} p_\theta(\mathbf{x}, \mathbf{z}) = \sum_{\mathbf{z} \in \mathcal{Z}} \frac{q(\mathbf{z})}{q(\mathbf{z})} p_\theta(\mathbf{x}, \mathbf{z}) = \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_\theta(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right]$$

where q can be any distribution over \mathbf{z} . We then apply Monte Carlo:

- 1 Sample $\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(k)}$ from $q(\mathbf{z})$ (instead of uniformly at random!)
- 2 Approximate expectation with sample average

$$p_\theta(\mathbf{x}) \approx \frac{1}{k} \sum_{j=1}^k \frac{p_\theta(\mathbf{x}, \mathbf{z}^{(j)})}{q(\mathbf{z}^{(j)})}$$

What is a good choice for $q(\mathbf{z})$? Intuitively, choose likely completions. When $q(\mathbf{z})$ is “good”, this approximation works very well! There are two challenges:

- 1 How to extend this to approximate the marginal log-likelihood $\log p_\theta(\mathbf{x})$?
- 2 How to choose a good q ?

Approximating the Marginal Log-Likelihood

We've seen a way to approximate marginal probabilities with importance sampling:

$$p_{\theta}(\mathbf{x}) = \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right] \approx \frac{1}{k} \sum_{j=1}^k \frac{p_{\theta}(\mathbf{x}, \mathbf{z}^{(j)})}{q(\mathbf{z}^{(j)})}$$

Next, we will extend this approach to approximate the marginal log-likelihood:

$$\log \left(\sum_{\mathbf{z} \in \mathcal{Z}} p_{\theta}(\mathbf{x}, \mathbf{z}) \right) = \log \left(\sum_{\mathbf{z} \in \mathcal{Z}} \frac{q(\mathbf{z})}{q(\mathbf{z})} p_{\theta}(\mathbf{x}, \mathbf{z}) \right) = \log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right] \right)$$

It's clear that

$$\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\log \left(\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right) \right] \neq \log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right] \right)$$


Jensen's Inequality

We want to approximate the marginal log-likelihood:

$$\log \left(\sum_{\mathbf{z} \in \mathcal{Z}} p_{\theta}(\mathbf{x}, \mathbf{z}) \right) = \log \left(\sum_{\mathbf{z} \in \mathcal{Z}} \frac{q(\mathbf{z})}{q(\mathbf{z})} p_{\theta}(\mathbf{x}, \mathbf{z}) \right) = \log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right] \right)$$

- $\log(\cdot)$ is a concave function. $\log(px + (1-p)x') \geq p \log(x) + (1-p) \log(x')$.
- Idea: use Jensen Inequality (for concave functions)

$$\log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} [f(\mathbf{z})] \right) = \log \left(\sum_{\mathbf{z}} q(\mathbf{z}) f(\mathbf{z}) \right) \geq \sum_{\mathbf{z}} q(\mathbf{z}) \log f(\mathbf{z})$$



Evidence Lower Bound via Jensen's Inequality

We want to approximate the marginal log-likelihood:

$$\log \left(\sum_{\mathbf{z} \in \mathcal{Z}} p_{\theta}(\mathbf{x}, \mathbf{z}) \right) = \log \left(\sum_{\mathbf{z} \in \mathcal{Z}} \frac{q(\mathbf{z})}{q(\mathbf{z})} p_{\theta}(\mathbf{x}, \mathbf{z}) \right) = \log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right] \right)$$

- We will use Jensen's Inequality (for concave functions)

$$\log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} [f(\mathbf{z})] \right) = \log \left(\sum_{\mathbf{z}} q(\mathbf{z}) f(\mathbf{z}) \right) \geq \sum_{\mathbf{z}} q(\mathbf{z}) \log f(\mathbf{z})$$

Choosing $f(\mathbf{z}) = \frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})}$, we obtain:

$$\log \left(\mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right] \right) \geq \mathbb{E}_{\mathbf{z} \sim q(\mathbf{z})} \left[\log \left(\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right) \right]$$

This is called the Evidence Lower Bound (**ELBO**).

The Evidence Lower Bound (ELBO)

The **evidence lower bound** (ELBO) holds for any q :

$$\begin{aligned}\log p(\mathbf{x}; \theta) &\geq \sum_{\mathbf{z}} q(\mathbf{z}) \log \left(\frac{p_{\theta}(\mathbf{x}, \mathbf{z})}{q(\mathbf{z})} \right) \\&= \sum_{\mathbf{z}} q(\mathbf{z}) \log p_{\theta}(\mathbf{x}, \mathbf{z}) - \underbrace{\sum_{\mathbf{z}} q(\mathbf{z}) \log q(\mathbf{z})}_{\text{Entropy } H(q) \text{ of } q} \\&= \sum_{\mathbf{z}} q(\mathbf{z}) \log p_{\theta}(\mathbf{x}, \mathbf{z}) + H(q)\end{aligned}$$

Variational Inference: We will maximize the log-likelihood by maximizing the ELBO while automatically choosing a good q .

How to Choose a Good q ?

Suppose $q(\mathbf{z})$ is **any** probability distribution over the hidden variables. A little bit of algebra reveals

$$D_{KL}(q(\mathbf{z})\|p(\mathbf{z}|\mathbf{x};\theta)) = \log p(\mathbf{x};\theta) - \underbrace{\sum_{\mathbf{z}} q(\mathbf{z}) \log p(\mathbf{z}, \mathbf{x}; \theta)}_{\text{-ELBO}} - H(q)$$

Rearranging, we get that

$$\log p(\mathbf{x};\theta) = \text{ELBO} + D_{KL}(q(\mathbf{z})\|p(\mathbf{z}|\mathbf{x};\theta)).$$

- The closer $q(\mathbf{z})$ is to $p(\mathbf{z}|\mathbf{x};\theta)$, the closer the ELBO is to the true log-likelihood.
- Confirms our previous importance sampling intuition: we should choose likely completions.
- In practice, the posterior $p(\mathbf{z}|\mathbf{x};\theta)$ is intractable to compute.
- **Strategy:** Find an approximate q close to $p(\mathbf{z}|\mathbf{x};\theta)$ via optimization.

Variational Inference and Learning

We have shown that at any datapoint \mathbf{x} , we have

$$\log p(\mathbf{x}; \theta) = \text{ELBO}(\theta, q) + D_{KL}(q(\mathbf{z}) \| p(\mathbf{z} | \mathbf{x}; \theta)) \geq \text{ELBO}(\theta, q)$$

Variational inference produces a tight approximation for $\log p(\mathbf{x}; \theta)$ by finding a q that makes $\text{ELBO}(q)$ tight:


- 1 We choose $q(\mathbf{z}; \phi)$ to be a (tractable) probability distribution over \mathbf{z} parameterized by ϕ (variational parameters).
- 2 We find a good ϕ^* by maximizing the ELBO and making it tight:

$$\log p(\mathbf{x}^{(i)}; \theta) \geq \max_{\phi} \text{ELBO}(\theta, \phi)$$

Once we have a good $\text{ELBO}(\theta, \phi^*)$ that is tight around $\log p(\mathbf{x}; \theta)$ we can do interesting things like optimize θ (next lecture!).


Example: Variational Inference Over Gaussians

Variational Inference: Optimize q to approximate the intractable $p(\mathbf{z}|\mathbf{x}; \theta)$.



- Example: Suppose $q(\mathbf{z}; \phi)$ is a (tractable) probability distribution over the hidden variables parameterized by ϕ (variational parameters)
 - For example, a Gaussian with mean and covariance specified by ϕ
$$q(\mathbf{z}; \phi) = \mathcal{N}(\phi_1, \phi_2)$$
- Variational inference: pick ϕ so that $q(\mathbf{z}; \phi)$ is as close as possible to $p(\mathbf{z}|\mathbf{x}; \theta)$. In the figure, the posterior $p(\mathbf{z}|\mathbf{x}; \theta)$ (blue) is better approximated by $\mathcal{N}(2, 2)$ (orange) than $\mathcal{N}(-4, 0.75)$ (green)

Example: Optimizing Likelihood with Missing Data



- Assume $p(\mathbf{x}^{top}, \mathbf{x}^{bottom}; \theta)$ assigns high probability to images that look like digits. In this example, we assume $\mathbf{z} = \mathbf{x}^{top}$ are unobserved (latent)
- Suppose $q(\mathbf{x}^{top}; \phi)$ is a (tractable) probability distribution over the hidden variables (missing pixels in this example) \mathbf{x}^{top} parameterized by ϕ (variational parameters)

$$q(\mathbf{x}^{top}; \phi) = \prod_{\text{unobserved variables } \mathbf{x}_i^{top}} (\phi_i)^{\mathbf{x}_i^{top}} (1 - \phi_i)^{(1 - \mathbf{x}_i^{top})}$$

- Is $\phi_i = 0.5 \forall i$ a good approximation to the posterior $p(\mathbf{x}^{top} | \mathbf{x}^{bottom}; \theta)$? No
- Is $\phi_i = 1 \forall i$ a good approximation to the posterior $p(\mathbf{x}^{top} | \mathbf{x}^{bottom}; \theta)$? No
- Is $\phi_i \approx 1$ for pixels i corresponding to the top part of digit **9** a good approximation? Yes

Summary

- Pros of Latent Variable Models:
 - Easy to build flexible models
 - Suitable for unsupervised learning
- Cons: Latent Variable Models:
 - Not tractable to evaluate likelihoods
 - Not tractable to train via maximum-likelihood
- Computing $p(\mathbf{x}; \theta)$ for arbitrary \mathbf{x} is hard in latent variable models
 - Variational inference produces a tight lower bound on $\log p(\mathbf{x}; \theta)$:

$$\log p(\mathbf{x}; \theta) \geq \sum_{\mathbf{z}} q(\mathbf{z}) \log p_{\theta}(\mathbf{x}, \mathbf{z}) + H(q)$$

for some good $q(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$ found by optimization

- Still, this is a complex approximation that only holds at one \mathbf{x} .
- Next lecture, we will look at how to scale this procedure to large datasets of \mathbf{x} .